Boston House Price Prediction

Methodology and Workflow

Author: Guillermo Daniel Bonilla

Date: March 2, 2026

Program: MIT Applied AI and Data Science Program

Course: Machine Learning Elective

Table of Contents

- **1.** Executive Summary
- **2.** Business Problem Overview
  - 2.1 Business Problem Overview
  - 2.2 Analysis Explanation in Achieving Business Objective
- **3.** Solution Approach / Methodology
- **4.** Data Overview
  - 4.1 Dataset Structure
  - 4.2 Categorical Variables
  - 4.3 Numerical Variables
  - 4.4 Data Quality Assessment
  - 4.5 Summary Descriptive Statistics
- **5.** Exploratory Data Analysis (EDA)
  - 5.1 Univariate Analysis
  - 5.2 Bivariate Analysis — Correlation Analysis
  - 5.3 Key Scatterplot Findings
- **6.** Data Preprocessing

    - 6.1 Outlier Handling
    - 6.2 Log Transformation of Target Variable (MEDV)
    - 6.3 Categorical Variable Encoding
    - 6.4 Data Partitioning
    - 6.5 Multicollinearity Diagnostics (VIF Analysis)
- **7.** Linear Regression Analysis and Assumptions
    - 7.1 Evaluation Metrics Definitions
    - 7.2 Model 1 — Full Feature Model
    - 7.3 Model 2 — Refined Model (Preferred)
    - 7.4 Multiple Linear Regression Equation
    - 7.5 Four OLS Assumptions Verification
    - 7.6 Cross-Validation Results
- **8.** Conclusions and Recommendations
    - 8.1 Actionable Insights
    - 8.2 Significance of Techniques Used
    - 8.3 Strategic Recommendations
    - 8.4 Business Success Translation
    - 8.5 Key Takeaways
- **9.** Workflow Summary

1. Executive Summary

The Boston Housing dataset was collected in 1970 from the **Boston Standard Metropolitan Statistical Area (SMSA)**, containing **506 unique home observations** across **13 variables**. This project builds an **ordinary least squares (OLS) multiple linear regression model** to capture variation in Boston home prices across neighborhoods and vicinities. The analysis comprises exploratory data analysis, feature engineering, and thorough diagnostics for multicollinearity.

Exploratory data analysis revealed robust spatial and socioeconomic heterogeneity across Boston's housing market. Key independent variables include **LSTAT** (percentage of lower-status population), **RM** (average number of rooms per dwelling), and **DIS** (weighted distances to five Boston employment centers), which showed strong relationships with home values. Houses farther from Boston employment centers were more likely to have large-lot zoning (ZN). Houses in areas with higher index of accessibility to radial highways (RAD) usually had higher property-tax rates (TAX).

Outlier values in each independent numerical variable were retained because they reflect real-world conditions. A **log transformation** was applied to the target-dependent median value (MEDV) to address skewness and improve linearity. RAD was **one-hot encoded** to avoid absolute multicollinearity, while CHAS remained binary. VIF analysis was conducted to identify and remove superfluous dummy variables with infinite or undefined VIF values. The remaining independent variables with acceptable VIF values showed no multicollinearity.

The final OLS model (**model2**) has an **R-squared of approximately 0.70**, explaining about 70% of the variation in home prices. RMSE, MAE, and MAPE are small and stable; no evidence of overfitting was observed. The most influential predictors (key drivers) and their interpretations are as follows:

Key Driver	Description	Impact on Home Values
CRIM	Per capita crime rate	Higher crime rates → lower home values
CHAS	Charles River adjacency	Homes near Charles River → valued more
NOX	Nitric oxide concentration	Higher nitric oxide pollution → lower values
RM	Average rooms per dwelling	More rooms → higher values
DIS	Weighted distances to employment centers	Farther from employment centers → lower values
PTRATIO	Pupil-teacher ratio	Higher pupil-teacher ratios (overcrowded schools) → lower values
LSTAT	Percentage lower-status	Higher lower-status

Key Driver	Description	Impact on Home Values
	population	population percentage → lower values

The project provides **data-driven discernment** of factors affecting Boston home pricing and delivers a statistically robust OLS model designed for effective interpretation and governance. Stakeholders — including **real estate businesses, home buyers, property investors, and housebuilders** — can utilize the model for operational decisions, pricing strategy, investment dispositions, marketing emphasis, and risk evaluation.

2. Business Problem Overview

2.1 Business Problem Overview

Different organizations and stakeholders — including home buyers, property investors, and housebuilders — need to understand which factors determined house values in Boston suburbs, towns, and neighborhoods in 1970. Without a data-driven approach, stakeholders risk incorrectly allocating resources, mispricing properties, and incorporating policies that do not capture achievable economic and community standards.

The core problem is to **correctly categorize the key performance variables that deterministically determine housing prices in Boston**. Misidentifying these drivers can lead to suboptimal investment decisions, inaccurate property valuations, and ineffective urban planning policies. A rigorous, model-based approach is essential for translating raw housing data into actionable intelligence.

2.2 Analysis Explanation in Achieving Business Objective

The business objective is to **accurately predict housing prices in Boston** across different towns and suburbs using compiled locality features. It is critical to identify the key factors driving house prices in distinct Boston suburbs and towns.

The analysis transforms housing and neighborhood data into legitimate insights by applying the following methods:

- **Exploratory Data Analysis (EDA)** methods to understand distributions, correlations, and spatial patterns
- **Data preprocessing techniques** to address skewed distributions, categorical encoding, and multicollinearity
- **OLS multiple linear regression modeling** to quantify the relationship between neighborhood attributes and home values

The project assesses how each factor — **CRIM, ZN, INDUS, CHAS, NOX, RM, AGE, DIS, RAD, TAX, PTRATIO, and LSTAT** — affects the target median value of owner-occupied homes (MEDV). The analysis ensures the model is interpretable and valid after addressing skewed distributions, categorical encoding, and multicollinearity.

3. Solution Approach / Methodology

The project follows a **data-driven, multi-faceted methodology** using Python-based statistical modeling and computations. The approach emphasizes Python-powered statistical analyses to deliver comprehensive data analysis and insights for strategic interventions and recommendations.

The dataset was collected from the Boston SMSA in 1970; stakeholders can use the house-prediction model to accurately describe home prices during that year. The solution approach includes **descriptive analytics** and **correlation analysis** as foundational pillars.

The complete methodological steps are as follows:

1. **Data Cleaning:** Verify no missing values and no duplicate entries across the entire dataset.
2. **Log Transformation:** Apply a natural log transformation to the target MEDV variable to improve linearity and stabilize variance.

3. **One-Hot Encoding:** Encode the RAD independent variable into binary dummy columns to avoid imposing ordinal assumptions.
4. **VIF Diagnostics:** Use Variance Inflation Factor analysis to identify and remove redundant categorical dummy variables exhibiting infinite or excessively high VIF values.
5. **Train-Test Split:** Use an 80:20 train-test split to evaluate model performance and guard against overfitting.
6. **Feature Selection:** Remove non-statistically significant independent variables (p-value > 0.05) to produce a refined, parsimonious model.

The methodology yields a **stable, high-quality model** that explains variation in housing prices while providing infrastructure for key decision-making in real estate planning, pricing, and policy development.

4. Data Overview

4.1 Dataset Structure

The dataset contains **506 unique homes** (rows) and **13 columns**, sourced from the Boston Standard Metropolitan Statistical Area (SMSA) in 1970. The variables span a diverse set of socioeconomic, environmental, and structural housing characteristics.

4.2 Categorical Variables

Variable	Description	Values
CHAS	Charles River dummy variable	0 (not river-adjacent), 1 (river-adjacent)
RAD	Index of accessibility to radial highways	1, 2, 3, 4, 5, 6, 7, 8, 24

4.3 Numerical Variables

Variable	Full Name	Type
CRIM	Per capita crime rate by town	float64
ZN	Proportion of residential land zoned for lots over 25,000 sq ft	float64
INDUS	Proportion of non-retail business acres per town	float64
NOX	Nitric oxide concentration (parts per 10 million)	float64
RM	Average number of rooms per dwelling	float64
AGE	Proportion of owner-occupied units built before 1940	float64
DIS	Weighted distances to five Boston employment centers	float64
TAX	Full-value property-tax rate per $10,000	int64
PTRATIO	Pupil-teacher ratio by town	float64
LSTAT	Percentage of lower status of the population	float64
MEDV	Median value of owner-occupied homes in $1,000 (target)	float64

4.4 Data Quality Assessment

- **Missing Values:** None — zero missing values across all 13 columns
- **Duplicate Entries:** None — zero duplicate rows detected
- **Dataset Shape:** 506 rows × 13 columns

4.5 Summary Descriptive Statistics

Variable	Count	Mean	Std	Min	25%	50%	75%	Max
CRIM	506	3.61	8.60	0.01	0.08	0.26	3.68	88.98
ZN	506	11.36	23.32	0.00	0.00	0.00	12.50	100.00
INDUS	506	11.14	6.86	0.46	5.19	9.69	18.10	27.74
CHAS	506	0.07	0.25	0.00	0.00	0.00	0.00	1.00
NOX	506	0.55	0.12	0.38	0.45	0.54	0.62	0.87
RM	506	6.28	0.70	3.56	5.89	6.21	6.62	8.78
AGE	506	68.57	28.15	2.90	45.02	77.50	94.07	100.00
DIS	506	3.80	2.11	1.13	2.10	3.21	5.19	12.13
RAD	506	9.55	8.71	1.00	4.00	5.00	24.00	24.00
TAX	506	408.24	168.54	187.00	279.00	330.00	666.00	711.00
PTRATIO	506	18.46	2.16	12.60	17.40	19.05	20.20	22.00
LSTAT	506	12.65	7.14	1.73	6.95	11.36	16.96	37.97
MEDV	506	22.53	9.20	5.00	17.02	21.20	25.00	50.00

Key observations from descriptive statistics:

- **CRIM:** Average of 3.61 indicates lower crime exposure overall, but the maximum of 88.98 reveals extreme high-crime pockets.
- **ZN:** Average of 11.36 with a median of 0.00 — most towns have no large-lot zoning; those that do tend to have substantial proportions.
- **RM:** Average of 6.28 rooms per dwelling — most homes are spacious with 5 to 7 rooms.
- **MEDV:** Average of $22,530 — the typical home is modest; values range from $5,000 to $50,000.
- **CHAS:** Mean of 0.07 — only approximately 7% of homes are adjacent to the Charles River.

- **AGE:** Average of 68.57% — the majority of owner-occupied units were built before 1940.
- **DIS:** Average of 3.80 — most homes are within moderate distance of employment centers.

5. Exploratory Data Analysis (EDA)

5.1 Univariate Analysis

The univariate analysis employed **bar plots** for categorical variables (CHAS and RAD) and **combined histogram-and-boxplot** visualizations for numerical variables. This dual-visualization approach captures both distributional shape and outlier presence simultaneously.

5.1.1 Categorical Variables

CHAS (Charles River Dummy Variable):

- **93.1%** of homes are not near the Charles River (CHAS = 0)
- **6.9%** of homes are river-adjacent (CHAS = 1)
- River-adjacent homes may correlate with higher MEDV, suggesting a premium for waterfront proximity

RAD (Index of Accessibility to Radial Highways):

- Distribution is uneven — indices **24 (26.1%)**, **5 (22.7%)**, and **4 (21.7%)** comprise approximately 70% of all observations
- Higher RAD values could suggest lower MEDV due to noise, traffic congestion, and pollution effects
- The concentration at RAD = 24 warrants further investigation for its relationship with tax rates

5.1.2 Numerical Variables

Variable	Distribution Shape	Median	Mean	Outlier Range	Key Observations
CRIM	Heavily right-skewed	0.26	3.61	15–80	Majority of towns have low crime; extreme outliers reflect high-crime pockets
ZN	Heavily right-skewed	0.00	11.36	30–100	Most towns have no large-lot zoning; outliers represent suburban areas
INDUS	Moderately right-skewed	9.69	11.14	—	Moderate industrial presence; spread across towns
NOX	Moderately right-skewed	0.54	0.55	—	Moderate pollution levels; tightly concentrated around the mean
RM	Slightly right-skewed (near normal)	6.21	6.28	—	Most homes have 5–7 rooms; nearly symmetric distribution
AGE	Slightly right-tailed	77.50	68.57	—	Most homes are older; high concentration of pre-

Variable	Distribution Shape	Median	Mean	Outlier Range	Key Observations
					1940 housing stock
DIS	Significantly right-skewed	3.21	3.80	—	Most homes are closer to employment centers; long tail toward suburbs
TAX	Moderately right-skewed	330	408.24	—	Bimodal tendency with concentration at lower and upper ranges
PTRATIO	Clustered distribution	19.05	18.46	—	Values tightly clustered between 16 and 22; limited variability
LSTAT	Significantly right-skewed	11.36	12.65	—	Moderate lower-status population percentage; right tail extends to 38%
MEDV	Right-skewed (target)	21.20	22.53	38K–50K	Strong concentration in $15K–$25K range; outliers at upper end

5.2 Bivariate Analysis — Correlation Analysis

A **Pearson correlation heatmap** was computed for all numerical variables to identify linear relationships between predictors and the target variable, as well as among the predictors themselves. The most significant correlations are summarized below:

Variable Pair	Pearson r	Interpretation
RM vs. MEDV	0.70	More rooms → higher home value
LSTAT vs. MEDV	−0.74	Higher lower-status percentage → lower home value
PTRATIO vs. MEDV	−0.51	Higher pupil-teacher ratio → lower home value
NOX vs. MEDV	−0.43	Higher pollution → lower home value
CRIM vs. MEDV	−0.39	Higher crime → lower home value
RAD vs. TAX	0.91	Higher highway access → higher property tax
NOX vs. INDUS	0.76	More industrial areas → higher pollution
DIS vs. NOX	−0.77	Farther from employment centers → less pollution
DIS vs. INDUS	−0.71	Suburban areas → less industrial land
AGE vs. DIS	−0.75	Older homes → closer to employment centers
AGE vs. NOX	0.73	Older homes → higher pollution
ZN vs. DIS	0.66	Farther from city → more large-lot zoning
INDUS vs. TAX	0.72	More industrial → higher tax rates

Important Note — RAD-TAX Correlation

The high RAD-TAX correlation (r = 0.91) was investigated further. After removing high TAX outliers, the correlation dropped to **0.25**, indicating the original correlation was driven by outlier observations concentrated at RAD = 24 and TAX = 666–711. This finding confirms the outlier-driven nature of this association.

5.3 Key Scatterplot Findings

Scatterplots were generated for all significant variable pairs to visually inspect the nature and strength of their relationships:

- **AGE vs. DIS:** Clear negative curvilinear relationship — older neighborhoods cluster near employment centers, while newer developments spread outward.
- **RAD vs. TAX:** Distinct clustering at RAD = 24 with high TAX values (666–711), confirming the outlier-driven correlation.
- **INDUS vs. TAX:** Positive linear trend — towns with more industrial acreage face higher property-tax rates.
- **LSTAT vs. MEDV:** Strong negative curvilinear relationship — as the percentage of lower-status population increases, home values decline sharply, with curvature suggesting diminishing effects at higher LSTAT levels.
- **RM vs. MEDV:** Positive linear trend — homes with more rooms consistently command higher values, with variance increasing at higher room counts.
- **INDUS vs. NOX:** Positive relationship — industrial concentration directly associates with pollution levels, with some nonlinearity at high INDUS values.
- **AGE vs. NOX:** Positive trend — older neighborhoods experience higher pollution, consistent with proximity to industrial and urban cores.

- **DIS vs. NOX:** Strong negative curvilinear pattern — pollution drops sharply as distance from employment centers increases, flattening at greater distances.

6. Data Preprocessing

6.1 Outlier Handling

Decision: No outlier values were removed. All outliers are true, real, and meaningful, reflecting truthful socioeconomic characteristics of the Boston housing market. Removing outliers would introduce several negative consequences:

- **Distribution distortion:** Removing extreme values would artificially narrow the data range and misrepresent the true variation in Boston's housing landscape.
- **Artificial sanitization bias:** Eliminating real-world conditions would produce a model that reflects a sanitized, unrealistic version of the housing market.
- **Reduced model transferability:** A model trained on truncated data would underperform when encountering genuine extreme values in practice.
- **Loss of significant correlation patterns:** Many meaningful correlations (e.g., RAD-TAX, CRIM-MEDV) depend on the full range of observations.

Validation Evidence

A separate OLS model was built with capped outliers (CRIM, ZN, DIS, PTRATIO, LSTAT). This alternative model produced **significantly weaker results** with lower R-squared and higher RMSE, MAE, and MAPE — confirming that retaining outliers yields superior model performance.

Skewed independent variables (CRIM, INDUS, DIS, LSTAT) were kept on their original scale; no log transformation was applied to these predictors, in order to maintain direct coefficient interpretability.

6.2 Log Transformation of Target Variable (MEDV)

Because MEDV is a **skewed target-dependent variable**, a natural log transformation was applied:

$$\mathbf{MEDV_log = log(MEDV)}$$

Aspect	Before Transformation	After Transformation
Variance	Unstable variance across predicted values	Stabilized, consistent variance
Linearity	No clear linear relationships with independent variables	Independent variables add fixed amounts to the target
Distribution	Right-skewed, non-normal	Nearly normally distributed with little skew
Residuals	Non-normal, heteroscedastic	Normally distributed and constant

Rationale: The log transformation improved model performance and produced more trustworthy multiple linear regression assumptions, ensuring that OLS estimates are reliable and interpretable.

6.3 Categorical Variable Encoding

- **CHAS:** Already binary (0/1) — no additional encoding needed.
- **RAD:** One-hot encoded to prevent the model from assuming meaningful distance, ordinal meaning, or equal spacing between categories. One-hot encoding enables the model to learn independent effects for each accessibility index group and improves interpretability.

6.4 Data Partitioning

The data was structured for modeling through the following steps:

- Independent variables partitioned into **X**; dependent target into **Y**
- An intercept term (constant) was added to X for the OLS model
- An **80:20 train-test split** was chosen given the small dataset size (506 observations)

Partition	Rows	Columns
X_train	404	20
X_test	102	20

6.5 Multicollinearity Diagnostics (VIF Analysis)

Multicollinearity occurs when two or more independent variables are highly correlated, causing unstable coefficients, inflated variance, and inaccurate model results. The Variance Inflation Factor (VIF) was used to detect and address multicollinearity iteratively.

Iterative VIF Process

Iteration 1: VIF computed for all features. **RAD_24** exhibited a VIF of **15.63** (threshold: VIF > 10) → *Dropped from the model.*

Iteration 2: VIF recomputed after removing RAD_24. **TAX** exhibited a VIF of **6.06** (threshold: VIF > 5) → *Dropped from the model.*

Final Result: All remaining features have VIF < 5, confirming no multicollinearity.

Final VIF Values

Feature	VIF	Status
CRIM	1.64	✅ Acceptable

Feature	VIF	Status
ZN	2.61	✅ Acceptable
INDUS	3.48	✅ Acceptable
CHAS	1.11	✅ Acceptable
NOX	4.52	✅ Acceptable
RM	1.96	✅ Acceptable
AGE	3.31	✅ Acceptable
DIS	4.42	✅ Acceptable
PTRATIO	2.20	✅ Acceptable
LSTAT	3.03	✅ Acceptable
RAD_2	1.52	✅ Acceptable
RAD_3	1.78	✅ Acceptable
RAD_4	2.14	✅ Acceptable
RAD_5	2.75	✅ Acceptable
RAD_6	1.46	✅ Acceptable
RAD_7	1.38	✅ Acceptable
RAD_8	1.51	✅ Acceptable

Boolean RAD dummy variables were converted to integer types before model building to ensure compatibility with the OLS regression framework.

7. Linear Regression Analysis and Assumptions

7.1 Evaluation Metrics Definitions

Metric	Definition	Optimal Direction
RMSE	Root Mean Squared Error — the square root of the average of squared	As low as possible

Metric	Definition	Optimal Direction
	differences between predicted and actual values. Penalizes larger errors more heavily.	
MAE	Mean Absolute Error — the average of the absolute differences between predicted and actual values. Provides an intuitive measure of average prediction error.	As low as possible
MAPE	Mean Absolute Percentage Error — the average of the absolute percentage differences between predicted and actual values. Expresses error as a percentage for interpretability.	As low as possible
R-squared (R^2)	Coefficient of Determination — the proportion of variance in the dependent variable that is explained by the independent variables. Ranges from 0 to 1.	As close to 1 as possible
Adjusted R^2	Modified R-squared that adjusts for the number of predictors in the model, penalizing the inclusion of non-significant variables.	As close to 1 as possible
p-value	The probability of observing the data given that the null hypothesis is true. A p-value < 0.05 indicates statistical significance at the 95% confidence level.	Below 0.05 for significance

7.2 Model 1 (model1) — Full Feature Model

Model 1 was built using all available features after VIF-based variable removal. Key statistics:

- **Observations:** 404
- **Residual Degrees of Freedom:** 386
- **R^2:** 0.780
- **Adjusted R^2:** 0.770

Statistically significant variables ($p < 0.05$):

- CRIM ($p = 0.000$), CHAS ($p = 0.006$), NOX ($p = 0.000$), RM ($p = 0.000$), DIS ($p = 0.000$), PTRATIO ($p = 0.000$), LSTAT ($p = 0.000$)

Non-significant variables ($p > 0.05$):

- ZN ($p = 0.052$), INDUS ($p = 0.950$), AGE ($p = 0.516$), RAD_2 ($p = 0.810$), RAD_3 ($p = 0.230$), RAD_4 ($p = 0.178$), RAD_5 ($p = 0.556$), RAD_6 ($p = 0.097$), RAD_7 ($p = 0.269$), RAD_8 ($p = 0.852$)

Model 1 Performance

Dataset	RMSE	MAE	MAPE (%)	R^2
Train	0.1896	0.1344	4.6719	0.7182
Test	0.2055	0.1553	5.4412	0.6154

No overfitting was detected — the model generalizes well to unseen data with MAPE of approximately 5%, indicating strong prediction accuracy.

7.3 Model 2 (model2) — Refined Model (Preferred)

Model 2 was constructed after removing the 10 non-significant variables (ZN, INDUS, AGE, and all RAD dummy variables). Key statistics:

- **Observations:** 404
- **R^2:** 0.771
- **Adjusted R^2:** 0.767
- **F-statistic:** 190.4

- All remaining variables are **highly significant** ($p < 0.05$)

Model 2 Coefficient Table

Feature	Coefficient	Std Error	t-stat	p-value	Interpretation
const	4.1437	0.205	20.221	0.000	Baseline log(MEDV)
CRIM	−0.0087	0.001	−7.036	0.000	Higher crime → ~0.87% decrease in price
CHAS	0.1070	0.037	2.897	0.004	Charles River adjacency → ~10.7% premium
NOX	−0.7785	0.143	−5.463	0.000	Higher pollution → ~7.8% price decrease per 0.1 unit
RM	0.0911	0.018	4.929	0.000	Each additional room → ~9.1% price increase
DIS	−0.0418	0.007	−5.748	0.000	Farther from jobs → ~4.2% decrease per unit
PTRATIO	−0.0376	0.005	−7.413	0.000	Higher pupil-teacher ratio → ~3.8% decrease per unit
LSTAT	−0.0290	0.002	−13.535	0.000	Higher lower-

Feature	Coefficient	Std Error	t-stat	p-value	Interpretation
					status % → ~2.9% decrease per unit

Model 2 Performance

Dataset	RMSE	MAE	MAPE (%)	R^2
Train	0.1935	0.1380	4.7822	0.7029
Test	0.2058	0.1574	5.5155	0.6164

Why Model 2 Is the Preferred Model

- **Slightly higher test R^2** (0.6164 vs. 0.6154) — better generalization
- **No statistically weak variables** — every predictor contributes meaningfully
- **No distortion of signs and magnitudes** — coefficients align with economic intuition
- **Higher interpretability** — fewer predictors yield clearer explanations
- **Greater stability** — reduced risk of overfitting with a parsimonious specification
- **Economic alignment** — all coefficient directions match established housing economics theory

7.4 Multiple Linear Regression Equation

$$\log(MEDV) = 4.1437 - 0.0087 \times CRIM + 0.1070 \times CHAS - 0.7785 \times NOX + 0.0911 \times RM - 0.0418 \times DIS - 0.0376 \times PTRATIO - 0.0290 \times LSTAT$$

Predictor strength classification:

Category	Predictors	Rationale
Strongest Predictors	NOX, RM, LSTAT	Largest absolute coefficient magnitudes; highest t-statistics
Moderate Predictors	PTRATIO, DIS, CHAS	Moderate coefficient sizes; all highly significant
Weakest Predictor	CRIM	Smallest absolute coefficient; still statistically significant

7.5 Four OLS Assumptions Verification

Assumption 1: Mean of Residuals ≈ 0

- **Result:** -4.81×10^{-15} (essentially zero)
- **Status:** ✅ **Met** — the model is not systematically over- or under-predicting

Assumption 2: Homoscedasticity (Goldfeld-Quandt Test)

- **F-statistic:** 1.088
- **p-value:** 0.279
- p-value > 0.05 → fail to reject the null hypothesis of equal variance
- F-statistic close to 1 → residual variances are very similar across subgroups
- **Status:** ✅ **Met** — residuals are homoscedastic (constant variance)

Assumption 3: Linearity of Variables (Residual Plot)

- Residuals do not form a noticeable strong pattern
- Residuals are randomly and uniformly scattered along the x-axis
- **Status:** ✅ **Met** — the linearity condition is satisfied

Assumption 4: Normality of Error Terms (Q-Q Plot and Histogram)

- **Q-Q Plot:** Points stay close to the theoretical reference line throughout; only minor tail deviations observed
- **Residual Histogram:** Strict bell-shaped curve peaking near zero; symmetric distribution; no signs of non-normality
- **Status:** ✅ **Met** — residuals follow a normal distribution

Assumptions Summary

All four OLS assumptions are satisfied, confirming that the model's coefficient estimates are unbiased, efficient, and reliable for inference and prediction.

7.6 Cross-Validation Results

A **10-fold cross-validation** was applied to assess the model's forecasting stability and generalizability:

Metric	Value	Interpretation
R^2	0.729 (±0.239)	Model explains ~73% of variance across all folds
MSE	0.040 (±0.023)	Small average squared error; consistent across folds

The cross-validation results confirm **forecasting stability** and support the model's suitability for continuous planning and decision-making applications.

8. Conclusions and Recommendations

8.1 Actionable Insights

7. **Allocate resources to improving housing quality:** Prioritize renovations and building larger, multi-room houses. The RM driver confirms that each additional room yields approximately a 9.1% increase in home value.
8. **Support environmental and pollution-reduction efforts:** Identify high-NOX neighborhoods and implement eco-friendly buffers such as tree planting, green corridors, and traffic rerouting to reduce nitric oxide concentrations.
9. **Enhance public transit connectivity:** Invest in improved transportation links to employment hubs for homes farther from employment centers. The DIS driver shows that each unit increase in distance corresponds to a ~4.2% decrease in value.
10. **Apportion funds to lower student-teacher ratios:** Invest in education infrastructure in poor-performing schools to decrease pupil-teacher ratios. The PTRATIO driver demonstrates that overcrowded schools suppress home values by ~3.8% per unit increase.
11. **Promote community programs:** Develop initiatives that enhance socioeconomic conditions across towns and suburbs. The LSTAT driver shows that higher lower-status population percentages consistently reduce home values.
12. **Implement crime-reduction and neighborhood-safety measures:** Identify crime hotspots, improve street lighting, and establish community safety initiatives. The CRIM driver confirms that higher crime rates lower home values by ~0.87% per unit.
13. **Leverage the preferred-housing-amenity effect:** Construct recreational areas near bodies of water and promote waterfront development. The CHAS driver reveals a ~10.7% premium for Charles River adjacency.

8.2 Significance of Techniques Used

- **Regression diagnostics and cross-validation** ensured statistically sound, operationally reliable insights that stakeholders can trust.
- **The modeling process** translated complex business dynamics into clear, understandable relationships between neighborhood attributes and home values.
- **Feature significance testing** identified the variables with the strongest influence, enabling targeted resource allocation.
- **Multicollinearity checking** ensured stable, interpretable coefficient estimates free from collinear distortion.
- **Performance metrics** (R^2, RMSE, MAE, MAPE) provided transparency into model accuracy and durability.

8.3 Strategic Recommendations

#	Recommendation	Description
1	**Pricing Strategy**	Set data-driven listing prices based on measurable neighborhood attributes (RM, LSTAT, NOX, PTRATIO). Classify undervalued homes that fall below predicted price thresholds as investment opportunities.
2	**Investment and Renovation Decisions**	Use CHAS and RM coefficients to prioritize property features for maximum return. Calculate ROI for housing upgrades based on the model's coefficient-driven value predictions.
3	**Site Selection**	Simulate future value scenarios under zoning changes. Identify areas

#	Recommendation	Description
		with concentrations of negative driving coefficient factors as targets for community improvement.
4	**Urban Planning and Policy**	Evaluate interventions (improving school funding, lowering NOX pollution) based on negative driving coefficients. Quantify the expected home-value impact of each policy lever.
5	**Buyer Segmentation**	Group purchasers by preferences — room space, school quality, employment center proximity — to tailor marketing and sales approaches to distinct buyer profiles.
6	**Marketing Strategy**	Advertise properties with the strongest positive predictors (CHAS, RM). Predict neighborhood alignment with buyer preferences to optimize marketing campaigns.

8.4 Business Success Translation

- The model provides **quantitative evidence** that minimizes financial risk and improves predicted price accuracy for all stakeholders.
- It identifies which predictors are most critical, **avoiding misallocation** of time, money, and organizational resources.
- Enables **competitive advantage** through data-driven, evidence-based explanations that outperform intuition-based decision-making.
- Supports **long-term strategic planning**: portfolio diversification, development planning, and policy advocacy grounded in empirical findings.

- Enhances **stakeholder communication** through clear coefficient interpretations that translate statistical outputs into business language.

8.5 Key Takeaways

- **Increased home size and Charles River adjacency increase home values** — RM and CHAS are the primary positive drivers.
- **Nitric oxide pollution levels highly negate home values** — NOX carries the largest absolute negative coefficient.
- **Socioeconomic factors frame buyer demand** — buyers pay more near high-performing schools and in higher socioeconomic demographics.
- **Close connectedness to employment centers matters** — homes farther away lose value, reflecting commute-cost sensitivity.
- **Lowering crime rates increases home values** — safety is a consistent, statistically significant positive influence.

9. Workflow Summary

The following presents the comprehensive end-to-end workflow employed in this project, organized into seven sequential phases.

Phase 1: Data Acquisition and Understanding

14. **Load the Boston Housing dataset** — 506 observations and 13 variables sourced from the Boston SMSA (1970).
15. **Inspect dataset structure** — identify 2 categorical variables (CHAS, RAD) and 10 continuous numerical variables plus 1 target (MEDV).
16. **Verify data quality** — confirm zero missing values and zero duplicate entries across all columns.
17. **Compute summary descriptive statistics** for all variables (mean, std, min, quartiles, max).

Phase 2: Exploratory Data Analysis

18. **Univariate analysis** — generate bar plots for categorical variables; combined histogram-boxplot visualizations for numerical variables.
19. **Identify distribution characteristics** — document distribution shapes, skewness, central tendency, and outliers across all variables.
20. **Bivariate analysis** — compute the Pearson correlation heatmap for all numerical variables.
21. **Investigate significant correlations** ($|r| \geq 0.7$): RM-MEDV (0.70), LSTAT-MEDV (−0.74), RAD-TAX (0.91), NOX-INDUS (0.76), DIS-NOX (−0.77), AGE-DIS (−0.75), AGE-NOX (0.73).
22. **Validate RAD-TAX correlation** — confirm it is driven by outliers (drops from 0.91 to 0.25 after removing high TAX values).
23. **Generate scatterplots** for all significant variable pairs to visualize relationship patterns.

Phase 3: Data Preprocessing

24. **Outlier retention decision** — retain all outliers as real-world socioeconomic indicators; validate with comparative model test.
25. **Log-transform the target variable**: MEDV_log = log(MEDV) to address skewness and improve linearity.
26. **Verify log-transformed distribution** — confirm MEDV_log is approximately normally distributed.
27. **Encode categorical variables** — one-hot encode RAD; retain CHAS as binary (0/1).
28. **Partition data** into features (X) and target (Y); add intercept term to X for OLS compatibility.
29. **Split data 80:20** into training (404 rows) and testing (102 rows) sets.

Phase 4: Multicollinearity Diagnostics

30. **Compute VIF for all features** — Iteration 1: identify RAD_24 with VIF = 15.63 (exceeds threshold of 10).
31. **Drop RAD_24 and recompute VIF** — Iteration 2: identify TAX with VIF = 6.06 (exceeds threshold of 5).
32. **Drop TAX and verify** — confirm all remaining features have VIF < 5 (no multicollinearity).
33. **Convert boolean RAD dummy variables** to integer types for model compatibility.

Phase 5: Model Building and Evaluation

34. **Build Model 1 (model1)** — full feature OLS regression on training data with all remaining predictors.
35. **Evaluate model1**: $R^2 = 0.780$ (train); identify 7 significant vs. 10 non-significant predictors using p-value threshold of 0.05.
36. **Remove 10 non-significant variables** (ZN, INDUS, AGE, RAD_2 through RAD_8) to refine the model.
37. **Build Model 2 (model2)** — refined OLS regression with 7 significant predictors only.
38. **Evaluate model2**: $R^2 = 0.771$ (train), $R^2 = 0.616$ (test); MAPE ~5% — strong prediction accuracy.
39. **Select model2 as the preferred model** for superior interpretability, stability, and economic alignment.

Phase 6: Assumption Verification

40. **Verify Assumption 1:** Mean of residuals ≈ 0 → Confirmed (-4.81×10^{-15}).
41. **Verify Assumption 2:** Homoscedasticity via Goldfeld-Quandt test → Confirmed ($F = 1.088$, $p = 0.279$).

42. **Verify Assumption 3:** Linearity via residual plot → Confirmed (random scatter, no systematic pattern).
43. **Verify Assumption 4:** Normality via Q-Q plot and histogram → Confirmed (bell-shaped curve, close to reference line).

Phase 7: Model Validation and Interpretation

44. **Apply 10-fold cross-validation**: $R^2 = 0.729$ (±0.239), MSE = 0.040 (±0.023) — confirms forecasting stability.
45. **Extract and interpret model coefficients** for all 7 predictors — translate statistical outputs into business-relevant insights.
46. **Formulate the final multiple linear regression equation** for operational use and stakeholder communication.
47. **Derive actionable insights and strategic recommendations** for real estate businesses, investors, policymakers, and buyers.

This workflow represents a **rigorous, end-to-end data science methodology** that ensures the final model is statistically sound, interpretable, and actionable for real estate business decision-making. Each phase builds upon the previous one, creating a transparent and reproducible analytical pipeline that translates raw housing data into quantitative evidence for strategic planning, pricing optimization, investment analysis, and urban policy development.